

# *The ER and Relational Models*

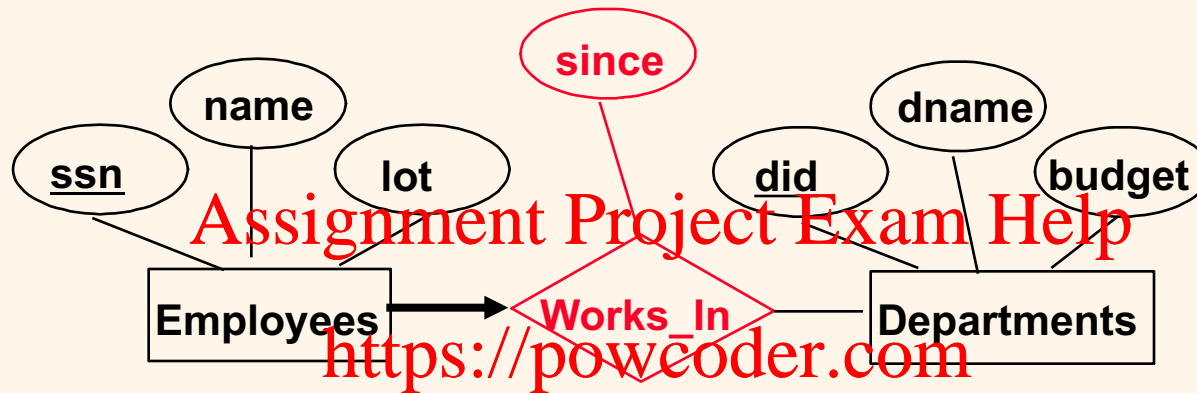
Assignment Project Exam Help

<https://powcoder.com>  
Module 3

Add WeChat powcoder

Professor Alex Brodsky  
Database Systems

# ER Review

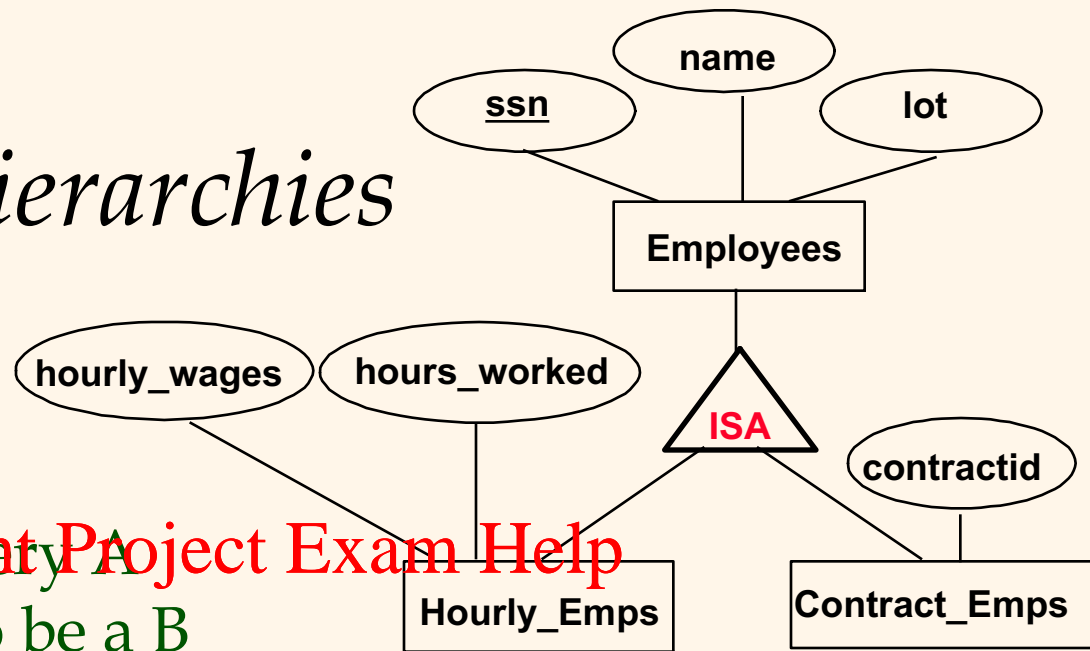


Add WeChat powcoder

Works\_in: m-to-1 relationship

Key for Works\_In: SSN (of employees)

# ISA ('is a') Hierarchies



❖ As in C++, or other PLs, attributes are inherited.

❖ If we declare **Assignment Project Exam Help** entity A, every A entity is also considered to be a B entity.

<https://powcoder.com>

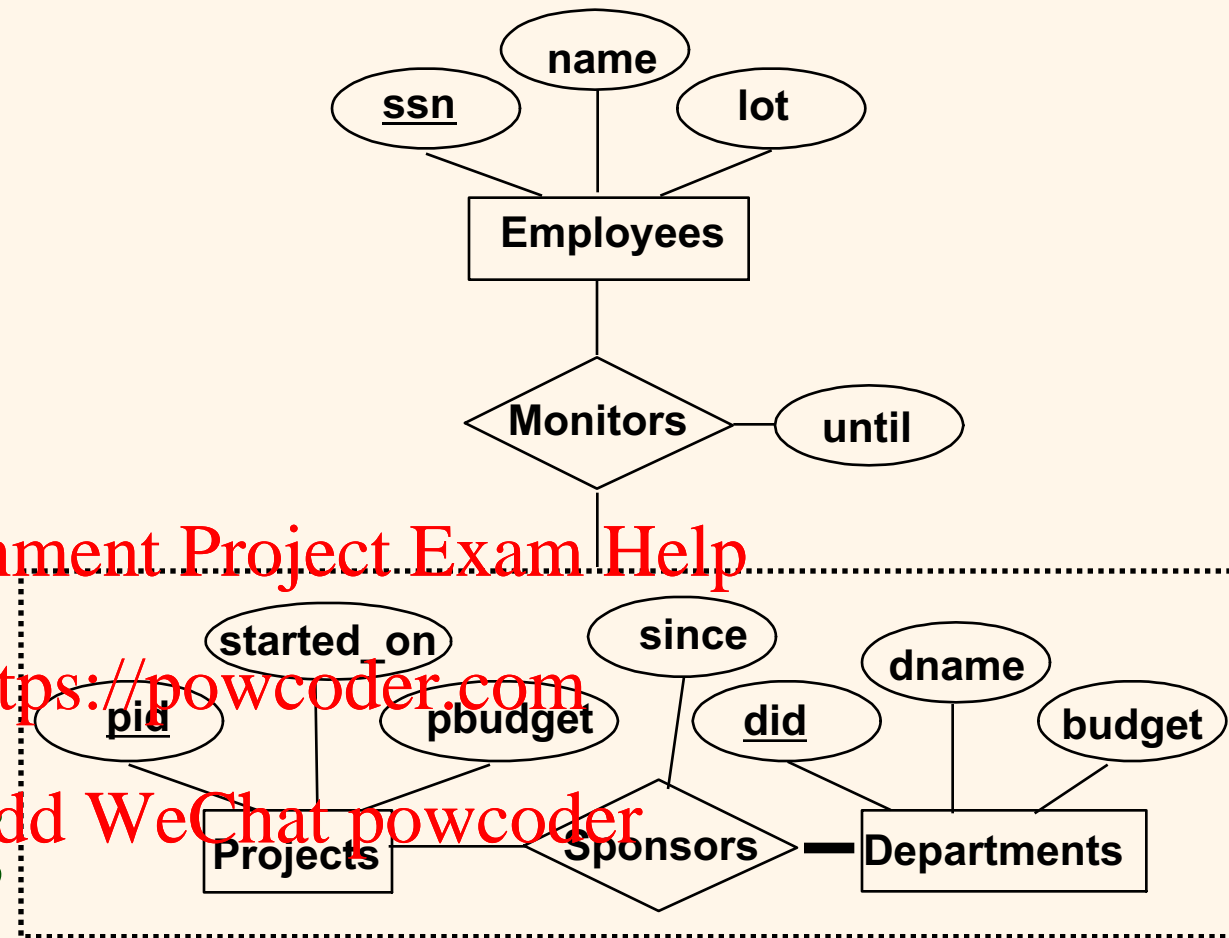
Add WeChat powcoder

- ❖ *Overlap constraints*: Can Joe be an Hourly\_Emps as well as a Contract\_Emps entity? *(Allowed/disallowed)*
- ❖ *Covering constraints*: Does every Employees entity also have to be an Hourly\_Emps or a Contract\_Emps entity? *(Yes/no)*
- ❖ Reasons for using ISA:
  - To add descriptive attributes specific to a subclass.
  - To identify entities that participate in a relationship.

# Aggregation

- ❖ Used when we have to model a relationship involving (entity sets and) a *relationship set*.

- Aggregation allows us to treat a relationship set as an entity set for purposes of participation in (other) relationships.



➡ *Aggregation vs. ternary relationship:*

- ❖ Monitors is a distinct relationship, with a descriptive attribute.
- ❖ Also, can say that each sponsorship is monitored by at most one employee.

# *Conceptual Design Using the ER Model*

## ❖ Design choices:

- Should a concept be modeled as an entity or an attribute?
- Should a concept be modeled as an entity or a relationship?
- Identifying relationships: Binary or ternary? Aggregation?

## ❖ Constraints in the ER Model:


- A lot of data semantics can (and should) be captured.
- But some constraints cannot be captured in ER diagrams.

# Entity vs. Attribute

- ❖ Should *address* be an attribute of Employees or an entity (connected to Employees by a relationship)?
- ❖ Depends upon the use we want to make of address information, and the semantics of the data:
  - ◆ If we have several addresses per employee, *address* must be an entity (since attributes cannot be set-valued).
  - ◆ If the structure (city, street, etc.) is important, e.g., we want to retrieve employees in a given city, *address* must be modeled as an entity (since attribute values are atomic).

# Entity vs. Attribute (Contd.)


- ❖ Works\_In2 does not allow an employee to work in a department for two or more periods.



- ❖ Similar to the problem of wanting to record


several addresses for an employee: we want to record *several values of the descriptive attributes for each instance of this relationship.*

<https://powcoder.com>  
Add WeChat powcoder



# Entity vs. Relationship

- ❖ First ER diagram OK if a manager gets a separate discretionary budget for each dept.

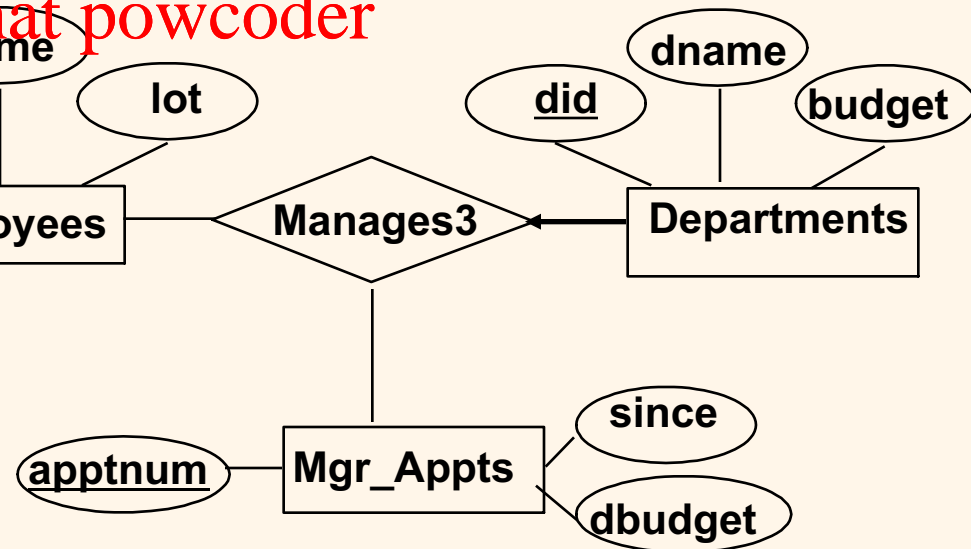


- ❖ What if a manager gets a discretionary budget that covers *all* managed depts?

<https://powcoder.com>  
Add WeChat powcoder

- **Redundancy** of *dbudget*, which is stored for each dept managed by the manager.

**Misleading:** suggests *dbudget* tied to managed dept.





# Binary vs. Ternary Relationships

- ❖ If each policy is owned by just 1 employee:

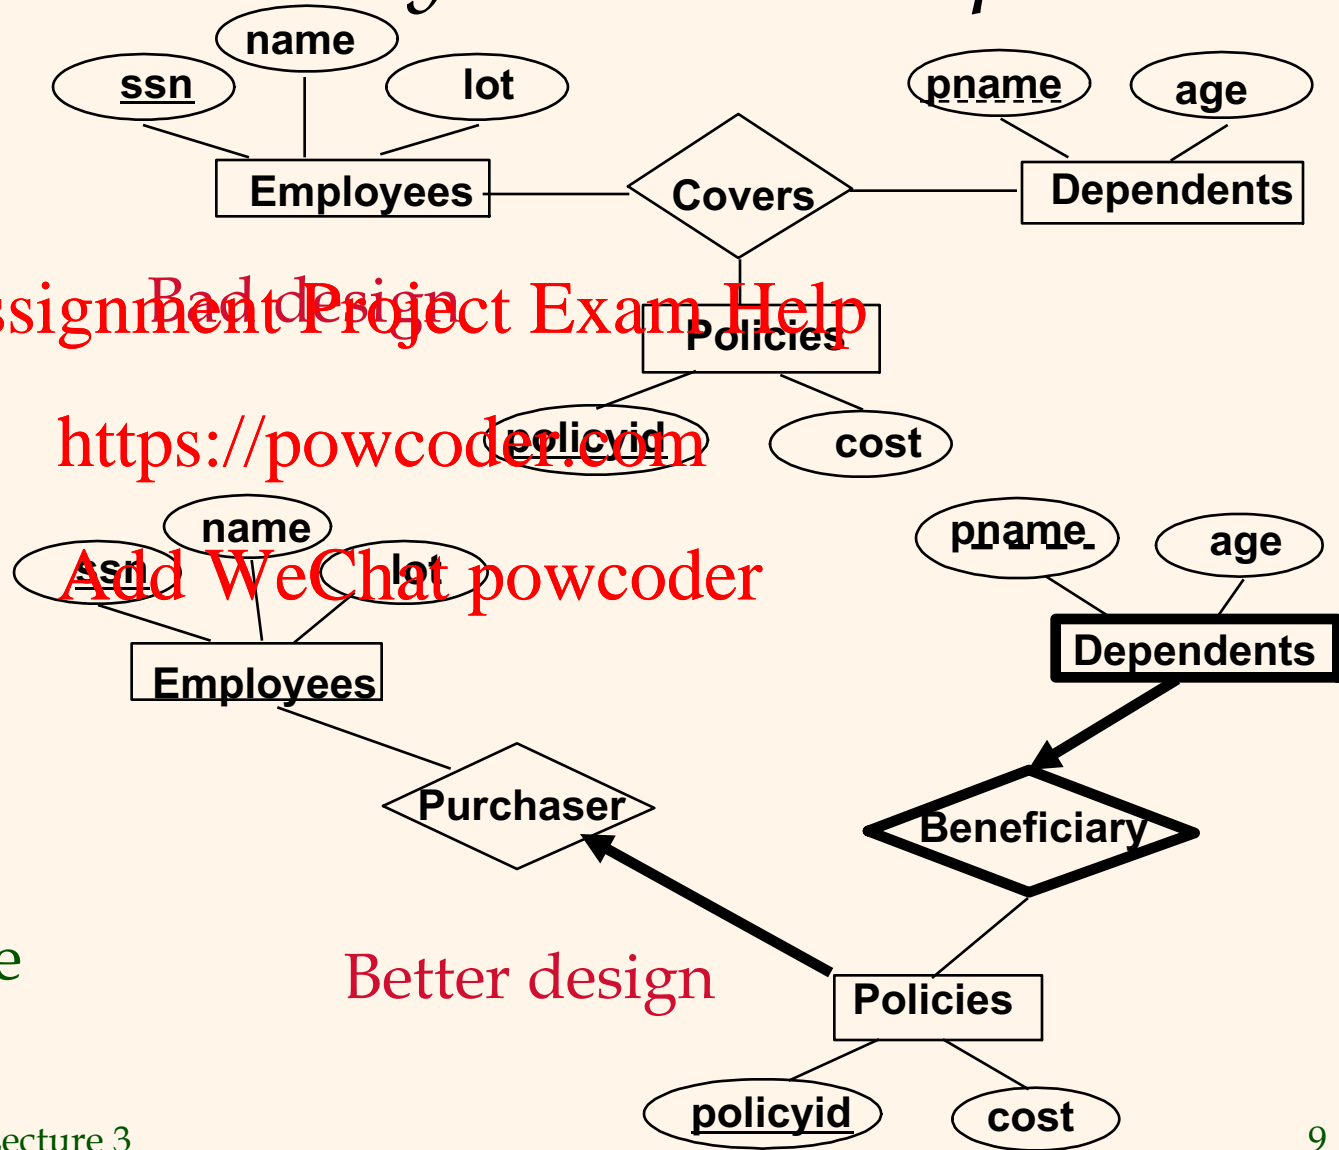
- Key constraint on Policies would mean policy can only cover 1 dependent!

- ❖ What are the additional constraints in the 2nd diagram?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Better design

# Binary vs. Ternary Relationships (Contd.)

- ❖ Previous example illustrated a case when two binary relationships were better than one ternary relationship
- ❖ An example in the other direction: a ternary relation *Contracts* relates entity sets *Parts*, *Departments* and *Suppliers* and has descriptive attribute *qty*. No combination of binary relationships is an adequate substitute:
  - S “can-supply” P, D “needs” P, and D “deals-with” S does not imply that D has agreed to buy P from S.
  - How do we record *qty*?

# Summary of Conceptual Design

- ❖ *Conceptual design follows requirements analysis,*
  - Yields a high-level description of data to be stored
- ❖ ER model popular for conceptual design
  - Constructs are expressive, close to the way people think about their applications.
- ❖ Basic constructs: *entities, relationships, and attributes* (of entities and relationships).
- ❖ Some additional constructs: *weak entities, ISA hierarchies, and aggregation.*
- ❖ Note: There are many variations on ER model.

# Summary of ER (Contd.)

- ❖ Several kinds of integrity constraints can be expressed in the ER model: *key constraints, participation constraints, and overlap/covering constraints* for ISA hierarchies. Some *foreign key constraints* are also implicit in the definition of a relationship set.
- Some constraints (notably, *functional dependencies*) cannot be expressed in the ER model.
  - Constraints play an important role in determining the best database design for an enterprise.

# Summary of ER (Contd.)

- ❖ ER design is *subjective*. There are often many ways to model a given scenario! Analyzing alternatives can be tricky, especially for a large enterprise. Common choices include:
  - Entity vs. attribute, entity vs. relationship, binary or n-ary relationship, whether or not to use ISA hierarchies, and whether or not to use aggregation.
- ❖ Ensuring good database design: resulting relational schema should be analyzed and refined further. FD information and normalization techniques are especially useful.

# Relational Model

- ❖ Relational Model = Structure + Operations
  - Structure: Relations (or Tables)
  - Operations: Relational Algebra, SQL.
- ❖ Most widely *implemented* model.
  - Vendors: IBM DB2, Microsoft SQL Server, Oracle, etc.
- ❖ Our design+implementation approach:
  - Step 1: ER design (ERD)
  - Step 2: Translate to Relational (Relational Schema)
  - Step 3: Querying over the relational model

# Relational Database: Definitions

- ❖ *Relational database*: a set of *relations*
- ❖ *Relation*: a set of *tuples*
  - *Instance*: a table, with rows and columns.  
#Rows = cardinality, #fields = degree / arity.
  - *Schema*: specifies name of relation, plus name and type of each column.
    - ◆ E.G. Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real).
- ❖ Can think of a relation as a *set* of rows or *tuples* (i.e., all rows are distinct).

## *Example Instance of Students Relation*

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eece	18	3.2
53650	Smith	smith@math	19	3.8

- ❖ Cardinality = 3, degree = 5, all rows distinct
- ❖ Do all columns in a relation instance have to be distinct?



# *Relational Query Languages*

- ❖ A major strength of the relational model: supports simple, powerful querying of data.
- ❖ Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.
  - The key: precise semantics for relational queries.
  - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.

# *The SQL Query Language*

- ❖ Developed by IBM (system R) in the 1970s
- ❖ Need for a standard since it is used by many vendors  
<https://powcoder.com>
- ❖ Standards:  
[Add WeChat powcoder](#)
  - SQL-86
  - SQL-89 (minor revision)
  - SQL-92 (major revision, current standard)
  - SQL-99 (major extensions)

# *The SQL Query Language*

- ❖ To find all 18 year old students, we can write:

Assignment Project Exam Help

SELECT *	sid	name	login	age	gpa
FROM Students S	53666	Jones	jones@cs	18	3.4
WHERE S.age=18	53688	Smith	smith@ee	18	3.2

<https://powcoder.com>  
Add WeChat powcoder

- To find just names and logins, replace the first line:

```
SELECT S.name, S.login
```

# Querying Multiple Relations

❖ What does the following query compute?

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade="A"
```

Given the following instance  
of Enrolled (is this possible if  
the DBMS ensures referential  
integrity?):

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

we get:

S.name	E.cid
Smith	Topology112

# *Creating Relations in SQL*

- ❖ Creates the Students relation. Observe that the type (domain) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.  

```
CREATE TABLE Students  
(sid: CHAR(20),  
name: CHAR(20),  
login: CHAR(10),  
age: INTEGER,  
gpa: REAL)
```
- ❖ As another example, the Enrolled table holds information about courses that students take.  

```
CREATE TABLE Enrolled  
(sid: CHAR(20),  
cid: CHAR(20),  
grade: CHAR(2))
```

# Destroying and Altering Relations

DROP TABLE Students

- ❖ Destroys the relation Students. The schema information and the tuples are deleted.

ALTER TABLE Students  
ADD COLUMN firstYear: integer

- ❖ The schema of Students is altered by adding a new field; every tuple in the current instance is extended with a *null* value in the new field.

# *Adding and Deleting Tuples*

- ❖ Can insert a single tuple using:

*Assignment Project Exam Help*  
`INSERT INTO Students (sid, name, login, age, gpa)  
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2)`

<https://powcoder.com>

- ❖ Can delete all tuples satisfying some condition (e.g., name = 'Smith'):

*Add WeChat powcoder*

```
DELETE  
FROM Students S  
WHERE S.name = 'Smith'
```

👉 *Powerful variants of these commands are available; more later!*

# Integrity Constraints (ICs)

- ❖ **IC:** condition that must be true for *any* instance of the database; e.g., domain constraints.
  - ICs are specified when schema is defined.
  - ICs are checked when relations are modified.
- ❖ A *legal* instance of a relation is one that satisfies all specified ICs.
  - DBMS should not allow illegal instances.
- ❖ If the DBMS checks ICs, stored data is more faithful to real-world meaning.
  - Avoids data entry errors, too!



# Primary Key Constraints

- ❖ A set of fields is a (candidate) key for a relation if:
  - 1. No two distinct tuples can have same values in all key fields, and
  - 2. This is not true for any subset of the key.
- Part 2 false? A *superkey*.
- If there's >1 candidate keys for a relation, one of the keys is chosen (by DBA) to be the *primary key*.
- ❖ E.g., *sid* is a key for Students. (What about *name*?) The set {*sid*, *gpa*} is a superkey.

# Primary and Candidate Keys in SQL

- ❖ Possibly many candidate keys (specified using **UNIQUE**), one of which is chosen as the *primary key*.
- ❖ “For a given student and course, there is a single grade.” vs. “Students can take only one course, and receive a single grade for that course; further, no two students in a course receive the same grade.”  
**Assignment Project Exam Help**  
**https://powcoder.com**  
**Add WeChat powcoder**  
**CREATE TABLE Enrolled**  
**(sid CHAR(20)**  
**cid CHAR(20),**  
**grade CHAR(2),**  
**PRIMARY KEY (sid,cid) )**
- ❖ Used carelessly, an IC can prevent the storage of database instances that arise in practice!  
**CREATE TABLE Enrolled**  
**(sid CHAR(20)**  
**cid CHAR(20),**  
**grade CHAR(2),**  
**PRIMARY KEY (sid),**  
**UNIQUE (cid, grade) )**

# Foreign Keys, Referential Integrity

- ❖ Foreign key : Set of fields in one relation that is used to `refer' to a tuple in another relation. (Must correspond to primary key of the second relation.) Like a `logical pointer'  
<https://powcoder.com>
- ❖ E.g. *sid* is a foreign key referring to **Students**:  
[Add WeChat powcoder](#)
  - Enrolled(*sid*: string, *cid*: string, *grade*: string)
  - If all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references.
  - Can you name a data model w/o referential integrity?
    - ◆ Links in HTML!

# Foreign Keys in SQL

- ❖ Only students listed in the Students relation should be allowed to enroll for courses.

**Assignment Project Exam Help**  
**<https://powcoder.com>**  
**Add WeChat powcoder**

```
CREATE TABLE Enrolled  
(sid CHAR(20), cid CHAR(20), grade CHAR(2),  
PRIMARY KEY (sid,cid),  
FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

# Enforcing Referential Integrity

- ❖ Consider Students and Enrolled; *sid* in Enrolled is a foreign key that references Students.
- ❖ What should be done if an Enrolled tuple with a non-existent student id is inserted? *(Reject it!)*
- ❖ What should be done if a Students tuple is deleted?
  - Also delete all Enrolled tuples that refer to it.
  - Disallow deletion of a Students tuple that is referred to.
  - Set *sid* in Enrolled tuples that refer to it to a *default sid*.
  - (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a special value *null*, denoting 'unknown' or 'inapplicable'.)
- ❖ Similar if primary key of Students tuple is updated.

# Referential Integrity in SQL/92

- ❖ SQL/92 supports all 4 options on deletes and updates.
  - Default is **NO ACTION** (*delete/update is rejected*)
  - **CASCADE** (also delete all tuples that refer to deleted tuple)
  - **SET NULL / SET DEFAULT** (sets foreign key value of referencing tuple)

CREATE TABLE Enrolled

(sid CHAR(20),

cid CHAR(20),

grade CHAR(2),

PRIMARY KEY (sid,cid),

FOREIGN KEY (sid)

REFERENCES Students

ON DELETE CASCADE

ON UPDATE SET DEFAULT )

# Where do ICs Come From?

- ❖ ICs are based upon the semantics of the real-world enterprise that is being described in the database relations.
- ❖ We can check a database instance to see if an IC is violated, but we can NEVER infer that an IC is true by looking at an instance.
  - An IC is a statement about *all possible* instances!
  - For example, we know *name* is not a key, but the assertion that *sid* is a key is given to us.
- ❖ Key and foreign key ICs are the most common; more general ICs supported too.


# Logical DB Design: ER to Relational

- ❖ Entity sets to tables.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



```
CREATE TABLE Employees
(ssn CHAR(11),
name CHAR(20),
lot INTEGER,
PRIMARY KEY (ssn))
```



# Relationship Sets to Tables

- ❖ In translating a relationship set to a relation, attributes of the relation must include:

- Keys for each participating entity set (as foreign keys)


- ◆ This set of attributes forms a *superkey* for the relation.

- All descriptive attributes.

```
CREATE TABLE Works_In(  
    ssn CHAR(11),  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY (ssn, did),  
    FOREIGN KEY (ssn)  
    REFERENCES Employees,  
    FOREIGN KEY (did)  
    REFERENCES Departments)
```


# Review: Key Constraints

- ❖ Each dept has at most one manager, according to the key constraint on **Manages**.




Assignment Project Exam Help  
<https://powcoder.com>


Add WeChat powcoder




1-to-1



1-to Many



Many-to-1



Many-to-Many

*Translation to relational model?*

# Translating ER Diagrams with Key Constraints

## ❖ Map relationship to a table:

- Note that **did** is the key now!
- Separate tables for Employees and Departments.

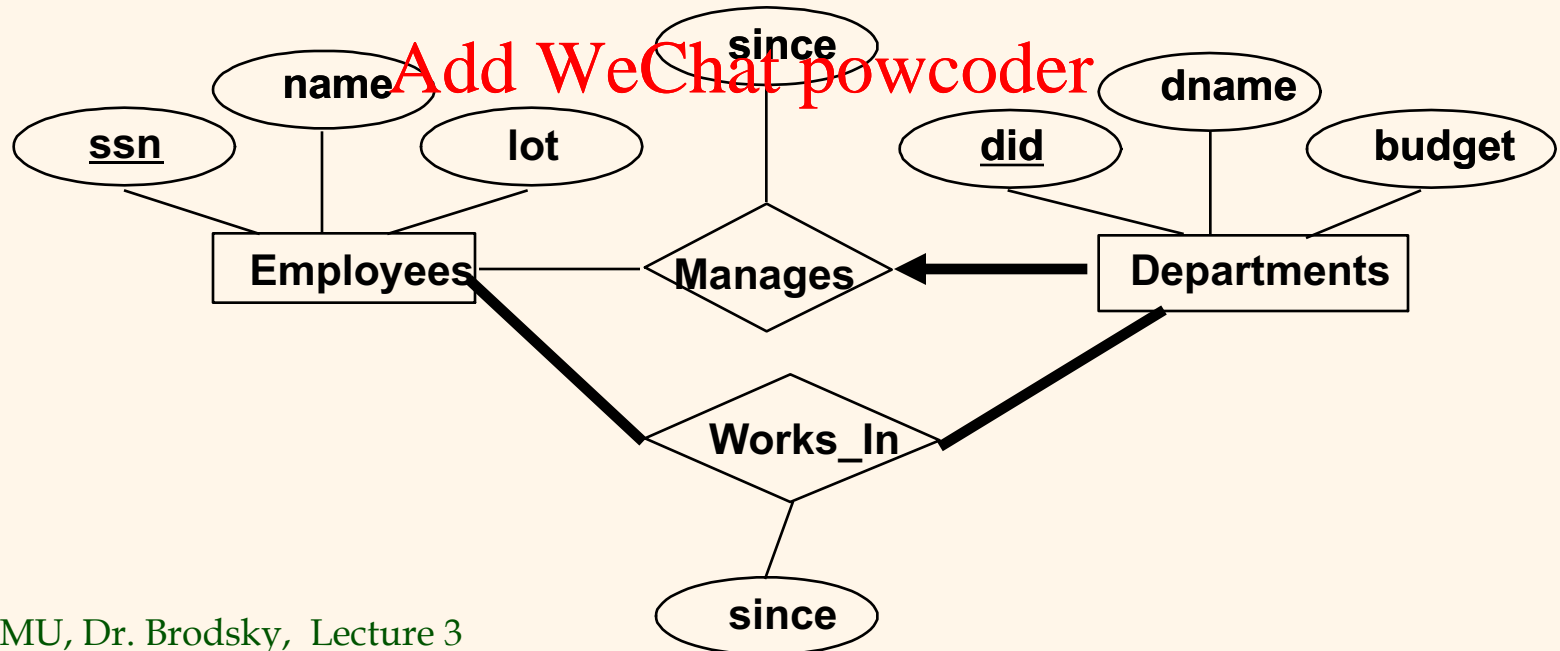
```
CREATE TABLE Manages(  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  FOREIGN KEY (did) REFERENCES Departments)
```

## ❖ Since each department has a unique manager, we could instead combine Manages and Departments.

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11),  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees)
```

# Review: Participation Constraints

- ❖ Does every department have a manager?
  - If so, this is a participation constraint: the participation of Departments in Manages is said to be *total* (vs. *partial*).
    - ◆ Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)



# *Participation Constraints in SQL*

- ❖ We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11) NOT NULL,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE NO ACTION)
```


# Review: Weak Entities

- ❖ A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
  - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
  - Weak entity set must have total participation in this identifying relationship set.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

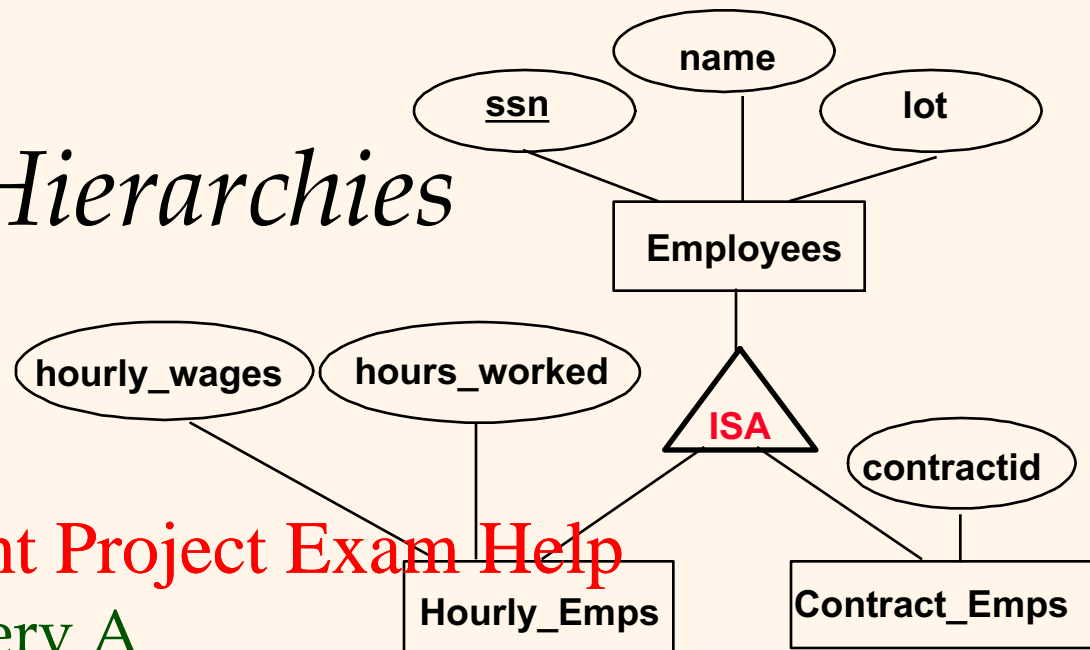


# Translating Weak Entity Sets

- ❖ Weak entity set and identifying relationship set are translated into a single table.
  - When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Dep_Policy(  
  pname CHAR(20),  
  age INTEGER,  
  cost REAL,  
  ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (pname, ssn),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE CASCADE)
```

# Review: ISA Hierarchies



❖ As in C++, or other PLs, attributes are inherited.

❖ If we declare A **ISA** B, every A entity is also considered to be a B entity.

**Add WeChat powcoder**

- ❖ *Overlap constraints*: Can Joe be an Hourly\_Emps as well as a Contract\_Emps entity? *(Allowed/disallowed)*
- ❖ *Covering constraints*: Does every Employees entity also have to be an Hourly\_Emps or a Contract\_Emps entity? *(Yes/no)*



# Translating ISA Hierarchies to Relations

## ❖ *General approach:*

- 3 relations: Employees, Hourly\_Emps and Contract\_Emps.

- ◆ *Hourly\_Emps*: Every employee is recorded in Employees. For hourly emps, extra info recorded in Hourly\_Emps (*hourly\_wages*, *hours\_worked*, *ssn*); must delete Hourly\_Emps tuple if referenced Employees tuple is deleted).

- ◆ Queries involving all employees easy, those involving just Hourly\_Emps require a join to get some attributes.

## ❖ *Alternative: Just Hourly\_Emps and Contract\_Emps.*

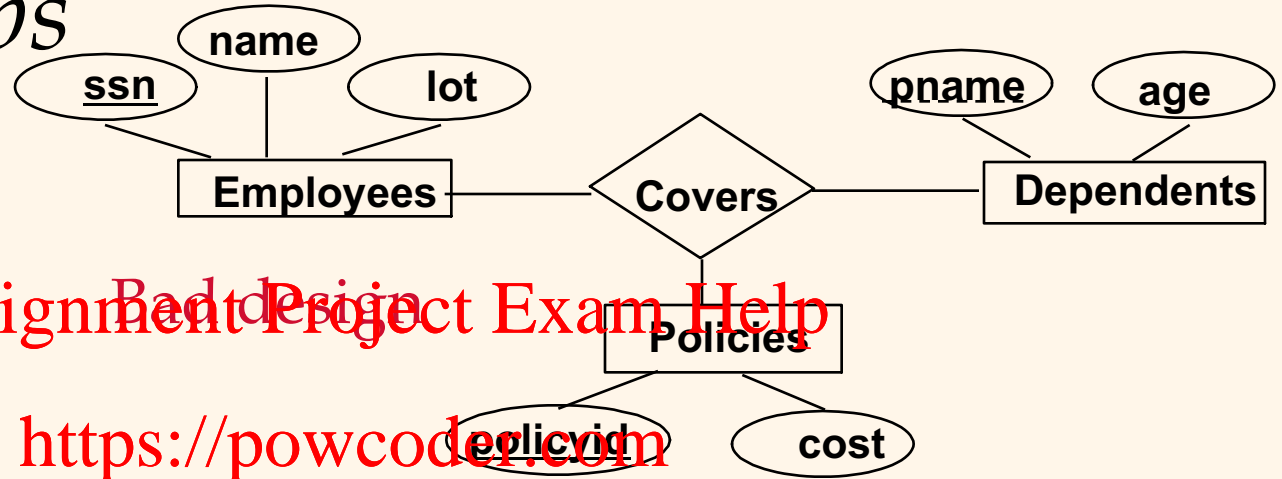
- *Hourly\_Emps*: *ssn*, *name*, *lot*, *hourly\_wages*, *hours\_worked*.
- Each employee must be in one of these two subclasses.

# Review: Binary vs. Ternary Relationships

- ❖ If each policy is owned by just 1 employee:

- Key constraint on Policies would mean policy can only cover 1 dependent!

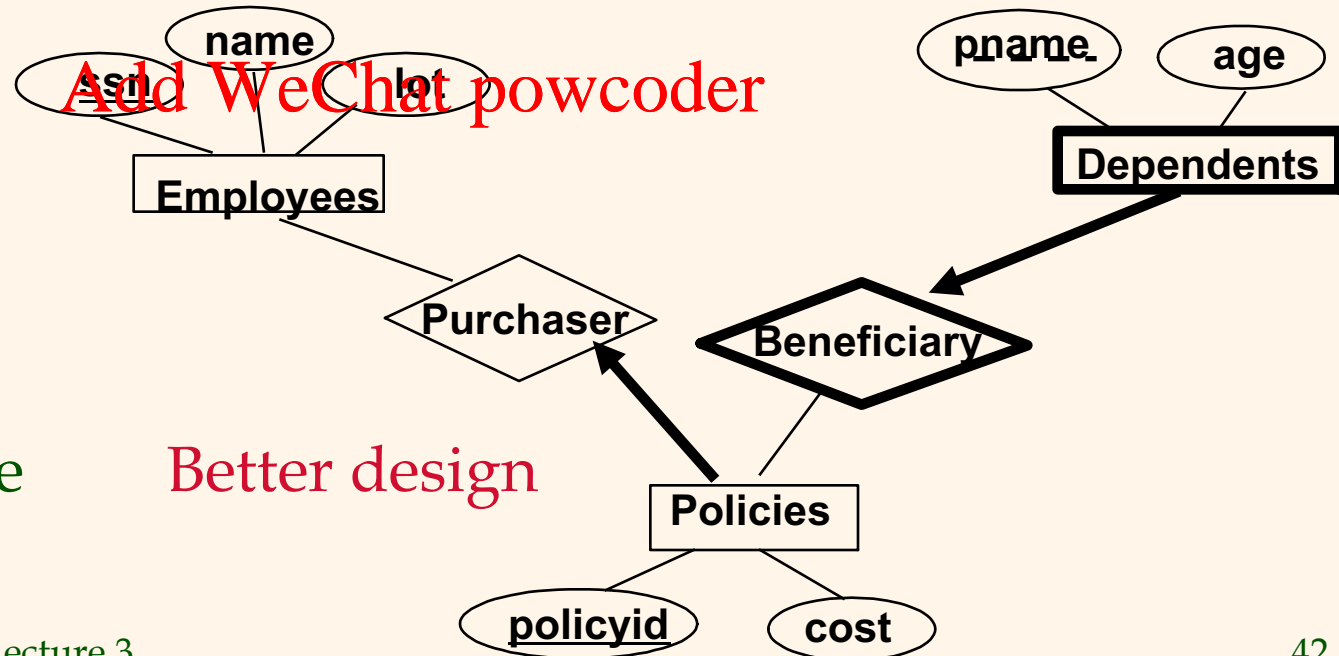
- ❖ What are the additional constraints in the 2nd diagram?



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Better design

# Binary vs. Ternary Relationships (Contd.)

- ❖ The key constraints allow us to combine Purchaser with Policies and Beneficiary with Dependents.

```
CREATE TABLE Policies (
```

```
    policyid INTEGER,
```

```
    cost REAL,
```

```
    ssn CHAR(11) NOT NULL,
```

```
    PRIMARY KEY (policyid).
```

```
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE CASCADE)
```

```
CREATE TABLE Dependents (
```

```
    pname CHAR(20),
```

```
    age INTEGER,
```

```
    policyid INTEGER,
```

```
    PRIMARY KEY (pname, policyid).
```

```
    FOREIGN KEY (policyid) REFERENCES Policies,  
    ON DELETE CASCADE)
```

- ❖ Participation constraints lead to **NOT NULL** constraints.
- ❖ What if Policies is a weak entity set?

# Relational Model: Summary

- ❖ A tabular representation of data.
- ❖ Simple and intuitive, currently the most widely used.
- ❖ Integrity constraints can be specified by the DBA, based on application semantics. DBMS checks for violations.
  - Two important ICs: primary and foreign keys
  - In addition, we *always* have domain constraints.
- ❖ Powerful and natural query languages exist.
- ❖ Rules to translate ER to relational model