

Relational Algebra

Assignment Project Exam Help

Module 4
<https://powcoder.com>

Add WeChat powcoder

Professor Alex Brodsky

Database Systems

Relational Query Languages

- ❖ Query languages: Allow manipulation and retrieval of data from a database.
- ❖ Relational model supports simple, powerful QLs:
 - Strong formal foundation based on logic.
 - Allows for much optimization.
- ❖ Query Languages != programming languages!
 - QLs not expected to be “Turing complete”.
 - QLs not intended to be used for complex calculations.
 - QLs support easy, efficient access to large data sets.

Formal Relational Query Languages

Two mathematical Query Languages form the basis for “real” languages (e.g. SQL), and for implementation:

- ① Relational Algebra: More operational, very useful for representing execution plans.
- ② Relational Calculus: Lets users describe what they want, rather than how to compute it.
(Non-operational, declarative.)

👉 Understanding Algebra is key to understanding SQL,
👉 and query processing!

Algebra Preliminaries

- ❖ A query is applied to *relation instances*, and the result of a query is also a relation instance.
 - *Schemas* of input relations for a query are *fixed* (but query will run regardless of instance!)
<https://powcoder.com>
 - The *schema* for the *result* of a given query is also *fixed*! Determined by definition of query language constructs.
Add WeChat powcoder

Example Instances

- ❖ “Sailors” and “Reserves” relations for our examples.

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

<https://powcoder.com>

Add WeChat powcoder

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Algebra Operations

- ❖ Look what we want to get from the following table:

Assignment Project Exam Help

<u>sid</u>	<u>sname</u>	<u>rating</u>	<u>age</u>
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Projection

sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

- ❖ Deletes attributes that are not in *projection list*.

- ❖ *Schema* of result contains exactly the fields in the projection list, with the same names that they had in the (only) input relation.

- ❖ Projection operator has to eliminate *duplicates*! (Why??)

- Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it. (Why not?)

$\pi_{sname, rating}(S2)$

age
35.0
55.5

$\pi_{age}(S2)$

Selection

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

❖ Selects rows that satisfy *selection condition*.

❖ No duplicates in result! (Why?)

❖ *Schema* of result identical to schema (only) input relation.

<https://powcoder.com>

$\sigma_{rating > 8}(S2) =$

sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

Composition of Operations

- ❖ *Result* relation can be the *input* for another relational algebra operation! (Operator composition.)

Assignment Project Exam Help

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<https://powcoder.com>

Add WeChat powcoder

$$\pi_{sname, rating}(\sigma_{rating > 8}(S2)) =$$

sname	rating
yuppy	9
rusty	10

What do we want to get from two relations?

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
58	rusty	10	35.0

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

What about: Who reserved boat 101?

Or: Find the name of the sailor who reserved boat 101.

Cross-Product

- ❖ Each row of S1 is paired with each row of R1.
- ❖ *Result schema* has one field per field of S1 and R1, with field names inherited.

Assignment Project Exam Help

sid1	sname	rating	age	sid2	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

➡ Renaming operator (because naming conflict):

$$\rho(sid \rightarrow sid1, S1) \times \rho(sid \rightarrow sid2, R1)$$

Why does this cross product help

Query: Find the name of the sailor who reserved boat 101.

Assignment Project Exam Help

$$CP = \rho(sid \rightarrow sid1, S1) \times \rho(sid \rightarrow sid2, R1)$$

$$Result = \pi_{Sname} (\sigma_{sid1=sid2 \text{ and } bid=101} (CP))$$

* Note my use of “temporary” relation CP.

Another example

- ❖ Find the name of the sailor having the highest rating.

Assignment Project Exam Help

$AllR = \pi_{ratingA}(\sigma_{rating > ratingA}(S2))$

Add WeChat powcoder

$Result? = \pi_{Sname}(\sigma_{rating < ratingA}(S2 \times AllR))$

What's in “Result?” ?

Does it answer our query?

Union, Intersection, Set-Difference

- ❖ All of these operations take two input relations, which must be union-compatible

- Same number of fields.
- 'Corresponding' fields have the same type.

- ❖ What is the *schema* of result?

sid	sname	rating	age
22	dustin	7	45.0

$S1 - S2$

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

$S1 \cup S2$

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$S1 \cap S2$

Back to our query

- ❖ Find the name of the sailor having the highest rating.

$AllR = \pi_{ratingA} (rating \rightarrow ratingA, S2)$

$Tmp = \pi_{Sid, Sname} (\sigma_{rating < ratingA} (S2 \times AllR))$

$Result = \pi_{Sname} (\pi_{Sid, Sname} (S2) - Tmp)$

- * Why not project on Sid **only** for Tmp?

Relational Algebra (Summary)

❖ Basic operations:

- Selection (σ) Selects a subset of rows from relation.
- Projection (π) Deletes unwanted columns from relation.
- Cross-product (\times) Allows us to combine two relations.
- Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
- Union (\cup) Tuples in reln. 1 and in reln. 2.
- Rename (ρ) Changes names of the attributes

❖ Since each operation returns a relation, **operations can be composed!** (Algebra is “closed”.)

❖ Use of temporary relations recommended.

Natural Join

- ❖ Natural-Join:

$S1 \bowtie R1$ Assignment Project Exam Help

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

- ❖ *Result schema* similar to cross-product, but only one copy of each field which appears in both relations.
- ❖ Natural Join = Cross Product + Selection on common fields + drop duplicate fields.

Query revisited using natural join

Query: Find the name of the sailor who reserved boat 101.

Assignment Project Exam Help

Result = $\pi_{Sname}(\sigma_{bid=101}(S1 \bowtie R1))$

Or Add WeChat powcoder

Result = $\pi_{Sname}(S1 \bowtie \sigma_{bid=101}(R1))$

Consider yet another query

- ❖ Find the sailor(s) who reserved all the red boats.

Assignment Project Exam Help

R1

B

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
22	103	10/11/96
58	102	11/12/96

<u>bid</u>	color
101	Red
102	Green
103	Red

<https://powcoder.com>

Add WeChat powcoder

Start an attempt

❖ who reserved what boat:

Assignment Project Exam Help

$S1 = \pi_{sid, bid}(R1) =$

Add WeChat powcoder

<u>sid</u>	<u>bid</u>
22	101
22	103
58	102

❖ All the red boats:

$S2 = \pi_{bid}(\sigma_{color=red}(B)) =$

<u>bid</u>
101
103

Division

- ❖ Not supported as a primitive operator, but useful for expressing queries like:

*Find sailors who have reserved **all** red boats.*

- ❖ Let $S1$ have 2 fields, x and y ; $S2$ have only field y :
 - $S1/S2 = \{ \langle x \rangle \mid \text{forall } \langle y \rangle \text{ in } R2 \text{ (there exists } \langle x, y \rangle \text{ in } R1) \}$
 - i.e., **$S1/S2$ contains all x tuples (sailors) such that for every y tuple (redboat) in $S2$, there is an xy tuple in $S1$ (i.e, x reserved y).**
- ❖ In general, x and y can be any lists of fields; y is the list of fields in $S2$, and $x \cup y$ is the list of fields of $S1$.

Examples of Division A/B

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

A

pno
p2

B1

sno
s1
s2
s3
s4

A/B1

pno
p2
p4

B2

sno
s1
s4

A/B2

pno
p1
p2
p4

B3

sno
s1

A/B3

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Find names of sailors who've reserved boat #103

❖ Solution 1: $\pi_{sname}((\sigma_{bid=103}(Reserves)) \bowtie Sailors)$

Assignment Project Exam Help

❖ Solution 2: $Temp1 = \sigma_{bid=103}(Reserves)$

<https://powcoder.com>
Add WeChat powcoder

$Temp2 = Temp1 \bowtie Sailors$

$Result = \pi_{sname}(Temp2)$

❖ Solution 3: $\pi_{sname}(\sigma_{bid=103}(Reserves \times Sailors))$

Find names of sailors who've reserved a red boat

- ❖ Information about boat color only available in Boats; so need an extra join

$\pi_{sname}((\sigma_{color='red'}Boats) \times Reserves \times Sailors)$

<https://powcoder.com>
Add WeChat powcoder

- ❖ A more efficient solution:

$\pi_{sname}(\pi_{sid}((\pi_{bid} \sigma_{color='red'}Boats) \times Res) \times Sailors)$

➡ *A query optimizer can find this given the first solution!*

Find sailors who've reserved a red or a green boat

- ❖ Can identify all red or green boats, then find sailors who've reserved one of these boats:

Assignment Project Exam Help

$Tempboats = \sigma_{color='red' \vee color='green'}(Boats)$

$Result = \pi_{sname}(Tempboats \bowtie Reserves \bowtie Sailors)$

- ❖ Can also define Tempboats using union! (How?)
- ❖ What happens if \vee is replaced by \wedge in this query?

Find sailors who've reserved a red and a green boat

- ❖ Previous approach won't work! Must identify sailors who've reserved red boats, sailors who've reserved green boats, then find the intersection (note that *sid* is a key for *Sailors*):

Tempred = $\pi_{sid}((\sigma_{color='red'}(Boats)) \bowtie Reserves)$

Tempgreen = $\pi_{sid}((\sigma_{color='green'}(Boats)) \bowtie Reserves)$

Result = $\pi_{sname}((Tempred \cap Tempgreen) \bowtie Sailors)$

Find the names of sailors who've reserved all boats

- ❖ Uses division; schemas of the input relations to / must be carefully chosen:

$$\text{Tempsids} = (\pi_{sid, bid}(\text{Reserves})) / (\pi_{bid}(\text{Boats}))$$

$$\text{Result} = \pi_{sname}(\text{Tempsids} \bowtie \text{Sailors})$$

- ❖ To find sailors who've reserved all 'Interlake' boats:

$$\dots / \pi_{bid}(\sigma_{bname='Interlake'}(\text{Boats}))$$

Summary

- ❖ The relational model has rigorously defined query languages that are simple and powerful.
- ❖ Relational algebra is more operational; useful as internal representation for query evaluation plans.
- ❖ Several ways of expressing a given query; a query optimizer should choose the most efficient version.