

Basic SQL

Assignment Project Exam Help

Module 6
<https://powcoder.com>

Add WeChat powcoder

Prof. Alex Brodsky

Database Systems

R1

Example Instances

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

- ❖ We will use these instances of the Sailors and Reserves relations in our examples.

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

- ❖ If the key for the Reserves relation contained only the attributes *sid* and *bid*, how would the semantics differ?

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Basic SQL Query

SELECT	<i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>

Assignment Project Exam Help

- ❖ *relation-list* A list of relation names (possibly with a *range-variable* after each name). One name can appear more than once, with different range-variable names.
<https://powcoder.com>
Add WeChat powcoder
- ❖ *target-list* A list of attributes of the relations in *relation-list*

Basic SQL Query

- ❖ qualification Comparisons (“Attr op const” or “Attr1 op Attr2”, where op is one of =, >, >=, <, <=, <>) combined using AND, OR and NOT.

<https://powcoder.com>
Add WeChat powcoder

Conceptual Evaluation Strategy

- ❖ Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
 - Compute the cross-product of *relation-list*.
 - Discard resulting tuples if they fail *qualifications*.
 - Delete attributes that are not in *target-list*.
- ❖ This strategy is probably the least efficient way to compute a query! An optimizer will find more efficient strategies to compute *the same answers*.

Set vs. Bag (Multiset)

- ❖ Set: {1, 2, 3}
 - No duplicates, no order
- ❖ Bag: {1, 2, 2, 2, 3, 3}
 - Duplicate possible, no order
- ❖ Membership test
 - Same
- ❖ SQL uses “Bag Semantics”
- ❖ Relational algebra uses “Set Semantics”

Conceptual Evaluation Strategy

Nested loops evaluation:

Assignment Project Exam Help Rn
<https://powcoder.com>
Add WeChat powcoder

```
SELECT    target-attribute-list
FROM      R1 R2 ... Rn
WHERE     qualification

Foreach tuple t1 in R1
...
Foreach tuple tn in Rn
  1. Substitute the attribute names in the qualification part
    with values from t1, ..., tn
  2. If the modified qualification part evaluates True
    then output target-attribute-values
    else do nothing
end
...
end
```

Example of Conceptual Evaluation

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid AND R.bid=103
```

Assignment Project Exam Help

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

A Note on Range Variables

- ❖ Really needed only if the same relation appears twice in the FROM clause. The previous query can also be written as:

<https://powcoder.com>
Add WeChat powcoder

```
SELECT S.sname  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid AND bid=103
```

OR

```
SELECT sname  
FROM Sailors, Reserves  
WHERE Sailors.sid=Reserves.sid  
AND bid=103
```

*It is good style,
however, to use
range variables
always!*

Find sailors who've reserved at least one boat

```
SELECT S.sid  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid
```

Assignment Project Exam Help

<https://powcoder.com>

❖ Is it different from algebra query below?

Add WeChat powcoder
 $\pi_{\text{Sid}}(\text{Sailor} \bowtie \text{Reserve})$

❖ How many times the same sid may appear?

❖ What is the effect of replacing *S.sid* by *S.sname* in the SELECT clause?

Expressions and Strings

```
SELECT S.age, age1=S.age-5, 2*S.age AS age2  
FROM Sailors S  
WHERE S.sname LIKE 'B_%B'
```

Assignment Project Exam Help

- ❖ Illustrates use of arithmetic expressions and string pattern matching. *Find triples (of ages of sailors and two fields defined by expressions) for sailors whose names begin and end with B and contain at least three characters.*
- ❖ **AS** and **=** are two ways to name fields in result.
- ❖ **LIKE** is used for string matching. **`_`** stands for any one character and **`%`** stands for 0 or more arbitrary characters.

Set Operations (back to set semantics)

❖ Set-ops:

- All duplicates removed!
- Union
- Intersect
- Except (difference)

❖ *Bag version*

- Union all

Select ...

From ...

Where ...

Set-op

Select ...

From ...

Where ...

Find sid's of sailors who've reserved a red or a green boat

- ❖ **UNION**: Can be used to compute the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).
Assignment Project Exam Help
<https://powcoder.com>
 - ❖ If we replace **OR** by **AND** in the first version, what do we get?
Add WeChat powcoder
 - ❖ Also available: **EXCEPT** (What do we get if we replace **UNION** by **EXCEPT**?)
- ```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
AND (B.color='red' OR B.color='green')
```
- ```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND
R.bid=B.bid
AND B.color='red'
UNION
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND
R.bid=B.bid
AND B.color='green'
```

Find sid's of sailors who've reserved a red and a green boat

```
SELECT S.sid
FROM Sailors S, Boats B1, Reserves R1,
      Boats B2, Reserves R2
WHERE S.sid=R1.sid AND R1.bid=B1.bid
      AND S.sid=R2.sid AND R2.bid=B2.bid
      AND (B1.color='red' AND B2.color='green')
```

- ❖ **INTERSECT**: Can be used to compute the intersection of any two *union compatible* sets of tuples.

- ❖ Included in the SQL/92 standard, but some systems don't support it.

- ❖ Contrast symmetry of the UNION and INTERSECT queries with how much the other versions differ.

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND
      R.bid=B.bid
      AND B.color='red'
```

INTERSECT

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND
      R.bid=B.bid
      AND B.color='green'
```

Qualification involving Sets

- ❖ Value **IN SET**
 - Value **NOT IN SET**
- ❖ **EXISTS SET**
 - **NOT EXISTS SET**
- ❖ **UNIQUE SET**
- ❖ Value θ **ANY SET**
- ❖ Value θ **ALL SET**
 - θ is one of =, >, >=, <, <=, <>
- ❖ Where do we get *SET* from?
 - *Select statement!*
 - *Nested queries.*

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Nested Queries

Find names of sailors who've reserved boat #103:

```
SELECT S.sname  
FROM Sailors S  
WHERE S.sid IN (SELECT R.sid  
                FROM Reserves R  
                WHERE R.bid=103)
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- ❖ A very powerful feature of SQL: a WHERE clause can itself contain an SQL query! (Actually, so can FROM and HAVING clauses.)
- ❖ To find sailors who've *not* reserved #103, use NOT IN.
- ❖ To understand semantics of nested queries, think of a nested loops evaluation: *For each Sailors tuple, check the qualification by computing the subquery.*

Two “equivalent” queries?

```
Select S.Sid
```

```
From Sailor S, Reserve R
```

```
Where S.Sid = R.Sid;
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

```
Select Sid
```

```
From Sailor
```

```
Where Sid in (Select Sid
```

```
From Reserve);
```

Nested Queries with Correlation

Find names of sailors who've reserved boat #103:

```
SELECT S.sname  
FROM Sailors S  
WHERE EXISTS (SELECT *  
               FROM Reserves R  
               WHERE R.bid=103 AND S.sid=R.sid)
```



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- ❖ **EXISTS** is another set comparison operator, like **IN**.
- ❖ If **UNIQUE** is used, and * is replaced by *R.bid*, finds sailors with exactly one reservation for boat #103. (**UNIQUE** checks for duplicate tuples; * denotes all attributes. Why do we have to replace * by *R.bid*?)
- ❖ Illustrates why, in general, subquery must be re-computed for each Sailors tuple.

More on Set-Comparison Operators

- ❖ Find sailors whose rating is greater than that of some sailor called Horatio:

Assignment Project Exam Help

<https://powcoder.com>
SELECT *
FROM Sailors S
WHERE S.rating > ANY (SELECT S2.rating
FROM Sailors S2
WHERE S2.sname='Horatio')

More

- ❖ Find sailors whose rating is greater than that of every other sailor.

Assignment Project Exam Help

<https://powcoder.com>
SELECT *
FROM Sailors S1
WHERE S1.rating > All (SELECT S2.rating
FROM Sailors S2
WHERE S2.sid <> S1.Sid);

Rewriting INTERSECT Queries Using IN

Find sid's of sailors who've reserved both a red and a green boat:

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
AND S.sid IN (SELECT S2.sid
FROM Sailors S2, Boats B2, Reserves R2
WHERE S2.sid=R2.sid AND R2.bid=B2.bid
AND B2.color='green')
```

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

- ❖ Similarly, EXCEPT queries re-written using NOT IN.
- ❖ To find *names* (not *sid's*) of Sailors who've reserved both red and green boats, just replace *S.sid* by *S.sname* in SELECT clause. (What about INTERSECT query?)

Division in SQL

Find sailors who've reserved all boats.

❖ Let's do it the hard way, without EXCEPT:

(1)

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
    ((SELECT B.bid
      FROM Boats B)
  EXCEPT
    (SELECT R.bid
      FROM Reserves R
      WHERE R.sid=S.sid))
```

(2) SELECT S.sname <https://powcoder.com>

FROM Sailors S

WHERE NOT EXISTS ((SELECT B.bid
FROM Boats B

Sailors S such that ... WHERE NOT EXISTS (SELECT R.bid
FROM Reserves R
there is no boat B without ... WHERE R.bid=B.bid
AND R.sid=S.sid))

a Reserves tuple showing S reserved B

Conclusions for now

- ❖ This is “core” part of SQL
 - ❖ Bag, nested loop
 - ❖ Similar to Relational Algebra, but not quite
 - ❖ A lot of extensions coming up!
- <https://powcoder.com>
Add WeChat powcoder