

# *More on SQL*

Assignment Project Exam Help

Module 7  
<https://powcoder.com>

Add WeChat powcoder

Prof. Alex Brodsky

Database Systems

R1

# Example Instances

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

- ❖ We will use these instances of the Sailors and Reserves relations in our examples.

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

- ❖ If the key for the Reserves relation contained only the attributes *sid* and *bid*, how would the semantics differ?

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

# Basic SQL Query

SELECT	<i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>

## Assignment Project Exam Help

- ❖ *relation-list* A list of relation names (possibly with a *range-variable* after each name). One name can appear more than once, with different range-variable names.  
<https://powcoder.com>  
Add WeChat powcoder
- ❖ *target-list* A list of attributes of the relations in *relation-list*

# Conceptual Evaluation Strategy

## Nested loops evaluation:

```
SELECT      target-attribute-list
FROM R1
WHERE qualification

Assignment Project Exam Help Rn
https://powcoder.com
Add WeChat powcoder

Foreach tuple t1 in R1
...
Foreach tuple tn in Rn
  1. Substitute the attribute names in the qualification part
    with values from t1, ..., tn
  2. If the modified qualification part evaluates True
    then output target-attribute-values
    else do nothing
end
...
end
```

# *Set Operations (back to set semantics)*

## ❖ Set-ops:

- All duplicates removed!
- Union
- Intersect
- Except (difference)

## ❖ *Bag version*

- Union all

Select ...

From ...

Where ...

*Set-op*

Select ...

From ...

Where ...

# Qualification involving Sets

- ❖ Value **IN SET**
  - Value **NOT IN SET**
- ❖ **EXISTS SET**
  - **NOT EXISTS SET**
- ❖ **UNIQUE SET**
- ❖ Value  $\theta$  **ANY SET**
- ❖ Value  $\theta$  **ALL SET**
  - $\theta$  is one of  $=, >, >=, <, <=, <>$
- ❖ Where do we get *SET* from?
  - *Select statement!*
  - *Nested queries.*

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Nested Queries

*Find names of sailors who've reserved boat #103:*

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN (SELECT R.sid
                FROM Reserves R
                WHERE R.bid=103)
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- ❖ A very powerful feature of SQL: a WHERE clause can itself contain an SQL query! (Actually, so can FROM and HAVING clauses.)
- ❖ To find sailors who've *not* reserved #103, use NOT IN.
- ❖ To understand semantics of nested queries, think of a nested loops evaluation: *For each Sailors tuple, check the qualification by computing the subquery.*

# *Conclusions for now*

- ❖ This is “core” part of SQL
  - ❖ Bag, nested loop
  - ❖ Similar to Relational Algebra, but not quite
  - ❖ A lot of extensions coming up!
- <https://powcoder.com>  
Add WeChat powcoder



# Post Processing

- ❖ Processing on the result of an SQL query:
  - Sorting: can sort the tuples in the output by any column (even the ones not appearing in the SELECT clause)
  - Duplicate removal
  - Example: <https://powcoder.com>

**Assignment Project Exam Help**  
**Add WeChat powcoder**  

```
SELECT Distinct S.sname  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid and R.bid=103  
Sort by S.sid asc, S.sname desc;
```

- ❖ *Aggregation operators*

# Aggregate Operators

- ❖ Significant extension of relational algebra.

COUNT (\*)  
COUNT ( [DISTINCT] A)  
SUM ( [DISTINCT] A)  
AVG ( [DISTINCT] A)  
MAX (A)  
MIN (A)

Assignment Project Exam Help *single column*

```
SELECT COUNT (*)  
FROM Sailors S
```

<https://powcoder.com>  
SELECT S.sname  
FROM Sailors S

```
SELECT AVG (S.age)  
FROM Sailors S  
WHERE S.rating=10
```

Add WeChat powcoder

```
WHERE S.rating = (SELECT MAX(S2.rating)  
FROM Sailors S2)
```

```
SELECT COUNT (DISTINCT S.rating)  
FROM Sailors S  
WHERE S.sname='Bob'
```

```
SELECT AVG ( DISTINCT S.age)  
FROM Sailors S  
WHERE S.rating=10
```

# *Find name and age of the oldest sailor(s)*

- ❖ The first query is illegal!  
(We'll look into the reason a bit later, when we discuss **GROUP BY**.)

```
SELECT S.sname, MAX (S.age)
FROM Sailors S
```

- ❖ The third query is equivalent to the second query, and is allowed in the SQL/92 standard, but is not supported in some systems.

```
SELECT S.sname, S.age
FROM Sailors S
WHERE S.age =
      (SELECT MAX (S2.age)
       FROM Sailors S2)
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE (SELECT MAX (S2.age)
       FROM Sailors S2)
      = S.age
```

# GROUP BY and HAVING

- ❖ So far, we've applied aggregate operators to all (qualifying) tuples. Sometimes, we want to apply them to each of several groups of tuples.
- ❖ Consider: *Find the age of the youngest sailor for each rating level.*
  - In general, we don't know how many rating levels exist, and what the rating values for these levels are!
  - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this (!):

For  $i = 1, 2, \dots, 10$ :

```
SELECT MIN (S.age)
FROM Sailors S
WHERE S.rating = i
```

# Queries With GROUP BY and HAVING

```
SELECT      [DISTINCT] target-list
FROM        relation-list
WHERE       qualification
GROUP BY    grouping-list
HAVING      group-qualification
```

Assignment Project Exam Help

<https://powcoder.com>

- ❖ The *target-list* contains (i) attribute names (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
  - The attribute list (i) must be a subset of *grouping-list*.  
Intuitively, each answer tuple corresponds to a *group*, and these attributes must have a single value per group. (A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.)

# Conceptual Evaluation

- ❖ The cross-product of *relation-list* is computed, tuples that fail *qualification* are discarded, 'unnecessary' fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
- ❖ The *group-qualification* is then applied to eliminate some groups. Expressions in *group-qualification* must have a single value per group!
  - In effect, an attribute in *group-qualification* that is not an argument of an aggregate op also appears in *grouping-list*. (SQL does not exploit primary key semantics here!)
- ❖ One answer tuple is generated per qualifying group.

Find the age of the youngest sailor with age  $\geq 18$ ,  
for each rating with at least 2 such sailors

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	Horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

- ❖ Only S.rating and S.age are mentioned in the SELECT, GROUP BY or HAVING clauses; other attributes 'unnecessary'.
- ❖ 2nd column of result is unnamed. (Use AS to name it.)

rating	age
1	33.0
7	45.0
7	35.0
8	55.5
10	35.0

rating	
7	35.0

Answer relation

*For each red boat, find the number of reservations for this boat*

```
SELECT B.bid, COUNT (*) AS scout  
FROM Sailors S, Boats B, Reserves R  
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```

<https://powcoder.com>

- ❖ Grouping over a join of three relations.
- ❖ What do we get if we remove *B.color='red'* from the WHERE clause and add a HAVING clause with this condition?
- ❖ What if we drop Sailors and the condition involving S.sid?



*Find the age of the youngest sailor with age > 18,  
for each rating with at least 2 sailors (of any age)*

```
SELECT S.rating, MIN (S.age)
```

```
FROM Sailors S
```

```
WHERE S.age > 18
```

```
GROUP BY S.rating
```

```
HAVING 1 < (SELECT COUNT (*)
```

```
FROM Sailors S2
```

```
WHERE S.rating=S2.rating)
```

- ❖ Shows HAVING clause can also contain a subquery.
- ❖ Compare this with the query where we considered only ratings with 2 sailors over 18!
- ❖ What if HAVING clause is replaced by:
  - HAVING COUNT(\*) >1

*Find those ratings for which the average age is the minimum over all ratings*

- ❖ Aggregate operations cannot be nested! **WRONG:**

```
SELECT S.rating
FROM Sailors S
WHERE S.age = (SELECT MIN (AVG (S2.age)) FROM Sailors S2)
```

**Assignment Project Exam Help**

<https://powcoder.com>

- ❖ Correct solution (in SQL/92):

**Add WeChat powcoder**

```
SELECT Temp.rating, Temp.avgage
FROM (SELECT S.rating, AVG (S.age) AS avgage
      FROM Sailors S
      GROUP BY S.rating) AS Temp
WHERE Temp.avgage = (SELECT MIN (Temp.avgage)
                    FROM Temp)
```

# *Continue from previous*

However, this should work on Oracle 8:

```
SELECT S.rating
FROM Sailors S
Group by S.rating
Having AVG(S.age) = (SELECT MIN (AVG (S2.age))
                     FROM Sailors S2
                     Group by rating);
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Conclusion

- ❖ Post processing on the result of queries is supported.
- ❖ Aggregation is the most complex “post processing”
  - “Group by” clause partition the results into groups
  - “Having” clause puts condition on groups (just like Where clause on tuples).

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder