

More on Normalization Theory & Database Programming

Assignment Project Exam Help

Module 11
<https://powcoder.com>

Add WeChat powcoder

Prof. Alex Brodsky

Database Systems

Example

A	B	C
1	1	2
1	1	3
2	2	3
2	2	2

FDs with A as the left side:	Satisfied by the relation?
$A \rightarrow A$	Yes (trivial FD)
$A \rightarrow B$	Yes
$A \rightarrow C$	No: tuples 1&2
$AB \rightarrow A$	Yes (trivial FD)
$AC \rightarrow B$	Yes

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat: powcoder

Example

Let $F = \{ A \rightarrow BC, B \rightarrow C \}$. Does $C \rightarrow AB$ in F^+ ?

Answer: No. Either of the following 2 reasons is ok:

Reason 1) $C^+ = C$, and does not include AB. <https://powcoder.com>

Reason 2) We can find a relation instance such that it satisfies F but does not satisfy $C \rightarrow AB$.

A	B	C
1	1	2
2	1	2

List all the non-trivial FDs in F^+

- ❖ Given $F = \{ A \rightarrow B, B \rightarrow C \}$. Compute F^+ (with attributes A, B, C).

	A	B	C	AB	AC	BC	ABC
A		√	√	√	√	√	√
B			√			√	
C							
AB			√			√	√
AC		√		√		√	√
BC							
ABC							

Attribute closure
$A^+ = ABC$
$B^+ = BC$
$C^+ = C$
$AB^+ = ABC$
$AC^+ = ABC$
$BC^+ = BC$
$ABC^+ = ABC$

Example

- ❖ Given $F = \{ A \rightarrow B, B \rightarrow C \}$. Find a relation that satisfies F :

A	B	C
1	1	2
2	1	2

Assignment Project Exam Help

<https://1powcoder.com>

Add WeChat powcoder

- ❖ Given $F = \{ A \rightarrow B, B \rightarrow C \}$. Find a relation that satisfies F but does not satisfy $B \rightarrow A$. Well, the above example suffices.
- ❖ Can you find an instance that satisfies F but not $A \rightarrow C$? No. Because $A \rightarrow C$ is in F^+

Examples

$R_1(A, B, C, D, E), A \rightarrow B, C \rightarrow D$

Candidate key: ACE.

Assignment Project Exam Help

Intuitively, B cannot be in a candidate key. Reason is that A is not determined by any other attributes (like E), and A has to be in a candidate key (because the a candidate key has to determine all the attributes).
Now if A is in a candidate key, B cannot be in the same candidate key, since we can drop B from the candidate without losing the property of being a “key”.

Same reasoning apply to others attributes.

Example

$R_1(A, B, C, D, E), A \rightarrow B, C \rightarrow D$ [Same as previous]

Which normal form?

Assignment Project Exam Help

Not in BCNF. This is the case where all attributes in the FDs appear in R_1 . We consider A , and C to see if either is a superkey or not. Obviously, neither A nor C is a superkey, and hence R_1 is not in BCNF. More precisely, we have $A \rightarrow B$ is in F^+ and non-trivial, but A is not a superkey of R_1 .

Example

$R_1(A, B, C, D, E), A \rightarrow B, C \rightarrow D$ [Same as previous]

Which normal form?

Assignment Project Exam Help

Not in 3NF. We have $A \rightarrow B$ in F^+ and non-trivial, but A is not a superkey of R_1 . Furthermore, B is not in any candidate key (since the only candidate key is ACE).

<https://powcoder.com>

Add WeChat powcoder

Example

- ❖ $R_2(A,B,F)$, $AC \rightarrow E$, $B \rightarrow F$.
- ❖ Candidate key: AB.
- ❖ BCNF? No, because $B \rightarrow F$.
- ❖ 3NF? No, because $B \rightarrow F$.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Example

- ❖ $R_4(D, C, H, G), A \rightarrow I, I \rightarrow A$
- ❖ Candidate key: DGHG
- ❖ BCNF? Yes
- ❖ 3NF? Yes

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Example

- ❖ $R(A,B,C,D,E,G,H)$, $F=\{AB \rightarrow C, AC \rightarrow B, B \rightarrow D, BC \rightarrow A, E \rightarrow G\}$
- ❖ Candidate keys?
 - H has to be in all candidate keys
 - E has to be in all candidate keys
 - G cannot be in any candidate key (since E is in all candidate keys already).
 - Since $AB \rightarrow C$, $AC \rightarrow B$ and $BC \rightarrow A$, we know no candidate key can have ABC together.
 - AEH, BEH, CEH are not superkeys.
 - Try ABEH, ACEH, BCEH. They are all superkeys. And we know they are all candidate keys (since above properties)
 - These are the only candidate keys: (1) each candidate key either contains A, or B, or C since no attributes other than A,B,C determine A, B, C, and (2) if a candidate key contains A, then it must contain either B, or C, and so on.

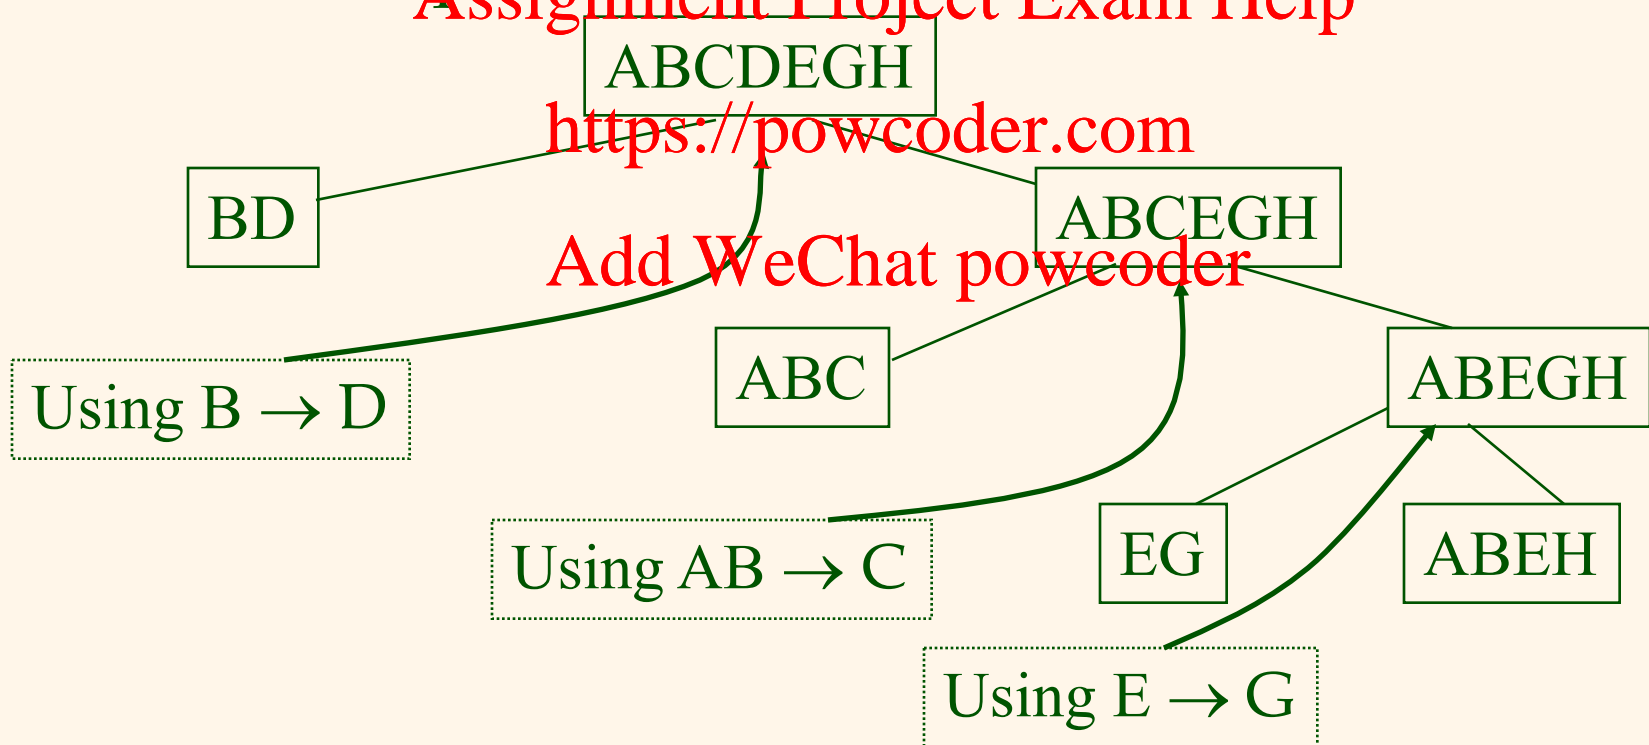
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Example

- ❖ Same as previous
- ❖ Not in BCNF, not in 3NF
- ❖ Decomposition:



Example

- ❖ $R(A,B,C,D,E,G,H)$, $F=\{AB \rightarrow C, AC \rightarrow B, B \rightarrow D, BC \rightarrow A, E \rightarrow G\}$
- ❖ Decomposition: $BD, ABC, EG, ABCH$
- ❖ Why good decomposition?
 - They are all in BCNF
 - Lossless-join decomposition
 - All dependencies are preserved.

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

Example

- ❖ $R=(ABDE)$ decomposed into $R1(ABD)$, $R2(ABE)$
- ❖ $F=\{AB \rightarrow DE\}$
- ❖ It is a dependency preserving decomposition!
 - $AB \rightarrow D$ can be checked in $R1$
 - $AB \rightarrow E$ can be checked in $R2$
 - $\{AB \rightarrow DE\}$ is equivalent to $\{AB \rightarrow D, AB \rightarrow E\}$

Database Programming



Embedded SQL

- ❖ SQL commands can be called from within a host language (e.g., C or COBOL) program.
 - SQL statements can refer to host variables (including special variables used to return status).
 - Must include a statement to *connect* to the right database.
- ❖ SQL relations are (multi-) sets of records, with no *a priori* bound on the number of records. No such data structure in C.
 - SQL supports a mechanism called a cursor to handle this.

Cursors

- ❖ Can declare a cursor on a relation or query statement (which generates a relation).
- ❖ Can *open* a cursor, and repeatedly *fetch* a tuple then *move* the cursor until all tuples have been retrieved.
 - Can use a special clause, called **ORDER BY**, in queries that are accessed through a cursor, to control the order in which tuples are returned.
 - ◆ Fields in **ORDER BY** clause must also appear in **SELECT** clause.
 - The **ORDER BY** clause, which orders answer tuples, is *only* allowed in the context of a cursor.
- ❖ Can also modify/delete tuple pointed to by a cursor.

Cursor that gets names of sailors who've reserved a red boat, in alphabetical order

```
EXEC SQL DECLARE sinfo CURSOR FOR  
SELECT S.sname  
FROM Sailors S, Boats B, Reserves R  
WHERE S.sid=R.sid AND R.bid=B.bid AND  
B.color='red'  
ORDER BY S.sname
```

Embedding SQL in C: An Example

```
char SQLSTATE[6];
EXEC SQL BEGIN DECLARE SECTION
char c_sname[20]; short c_minrating; float c_age;
EXEC SQL END DECLARE SECTION
c_minrating = random();
EXEC SQL DECLARE sinfo CURSOR FOR
    SELECT S.sname, S.age FROM Sailors S
    WHERE S.rating > :c_minrating
    ORDER BY S.sname;
do {
    EXEC SQL FETCH sinfo INTO :c_sname, :c_age;
    printf("%s is %d years old\n", c_sname, c_age);
} while (SQLSTATE != '02000');
EXEC SQL CLOSE sinfo;
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Database APIs: Alternative to embedding

Rather than modify compiler, add library with database calls (API)

- ❖ special standardized interface: procedures/ objects
- ❖ passes SQL strings from language, presents result sets in a language-friendly way
- ❖ Microsoft's ODBC becoming C/C++ standard on Windows
- ❖ Sun's JDBC a Java equivalent
- ❖ Supposedly DBMS-neutral
 - a “driver” traps the calls and translates them into DBMS-specific code
 - database can be across a network

SQL API in Java (JDBC)

```
Connection con = // connect
    DriverManager.getConnection(url, "login", "pass");
Statement stmt = con.createStatement(); // set up stmt
String query = "SELECT name, rating FROM Sailors";
ResultSet rs = stmt.executeQuery(query);
try { // handle exceptions
    // loop through result tuples
    while (rs.next()) {
        String s = rs.getString("name");
        Int n = rs.getFloat("rating");
        System.out.println(s + "    " + n);
    }
} catch(SQLException ex) {
    System.out.println(ex.getMessage ()
        + ex.getSQLState () + ex.getErrorCode ());
}
```

TRANSACTION MANAGEMENT

Airline Reservations

many updates

Statistical Abstract of the US

many queries

Atomicity – all or nothing principle

Serializability – the effect of transactions as if they occurred one at a time

Items – units of data to be controlled

fine-grained – small items

**course-grained – large items
(granularity)**

Controlling access by locks

Read – sharable with other readers shared

Write – not sharable with anyone else exclusive

Model – (item, locktype, transaction ID)

Transactions

Transaction = a unit of work that must be:

1. *Atomic* = either all work is done, or none of it.
2. *Consistent* = relationships among values maintained.
<https://powcoder.com>
3. *Isolated* = appear to have been executed when no other DB operations were being performed.
Add WeChat powcoder
 - Often called *serializable* behavior.
4. *Durable* = effects are permanent even if system crashes.

Commit/Abort Decision

Each transaction ends with either:

1. *Commit* = the work of the transaction is installed in the database; previously its changes may be invisible to other transactions.
 2. *Abort* = no changes by the transaction appear in the database; it is as if the transaction never occurred.
 - ROLLBACK is the term used in SQL and the Oracle system.
- ❖ In the ad-hoc query interface (e.g., PostgreSQL psql interface), transactions are single queries or modification statements.
- Oracle allows SET TRANSACTION READ ONLY to begin a multistatement transaction that doesn't change any data, but needs to see a consistent “snapshot” of the data.
- ❖ In program interfaces, transactions begin whenever the database is accessed, and end when either a COMMIT or ROLLBACK statement is executed.

Example

`Sells(bar, beer, price)`

- ❖ Joe's Bar sells Bud for \$2.50 and Miller for \$3.00.
- ❖ Sally is querying the database for the highest and lowest price Joe charges:

(1) `SELECT MAX(price) FROM Sells
WHERE bar = 'Joe's Bar';`

(2) `SELECT MIN(price) FROM Sells
WHERE bar = 'Joe's Bar';`

- ❖ At the same time, Joe has decided to replace Miller and Bud by Heineken at \$3.50:

(3) `DELETE FROM Sells
WHERE bar = 'Joe's Bar' AND
(beer = 'Miller' OR beer = 'Bud');`

(4) `INSERT INTO Sells
VALUES('Joe's bar', 'Heineken', 3.50);`

- ❖ If the order of statements is 1, 3, 4, 2, then it appears to Sally that Joe's minimum price is greater than his maximum price.
- ❖ Fix the problem by grouping Sally's two statements into one transaction, *e.g.*, with one SQL statement.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Example: Problem With Rollback

- ❖ Suppose Joe executes statement 4 (insert Heineken), but then, during the transaction thinks better of it and issues a ROLLBACK statement.
- ❖ If Sally is allowed to execute her statement 1 (find max) just before the rollback, she gets the answer \$3.50, even though Joe doesn't sell any beer for \$3.50.
- ❖ Fix by making statement 4 a transaction, or part of a transaction, so its effects cannot be seen by Sally unless there is a COMMIT action.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

SQL Isolation Levels

Isolation levels determine what a transaction is allowed to see.

The declaration, valid for one transaction, is:

```
SET TRANSACTION ISOLATION LEVEL X;
```

where: **Assignment Project Exam Help**

- ❖ **X = SERIALIZABLE**: this transaction must execute as if at a point in time, where all other transactions occurred either completely before or completely after.
https://powcoder.com
Add WeChat powcoder
 - Example: Suppose Sally's statements 1 and 2 are one transaction and Joe's statements 3 and 4 are another transaction. If Sally's transaction runs at isolation level SERIALIZABLE, she would see the `Sells` relation either before or after statements 3 and 4 ran, but not in the middle.

- ❖ $X = \text{READ COMMITTED}$: this transaction can read only committed data.
 - Example: if transactions are as above, Sally could see the original `Sells` for statement 1 and the completely changed `Sells` for statement 2.
- ❖ $X = \text{REPEATABLE READ}$: if a transaction reads data twice, then what it saw the first time, it will see the second time (it may see more the second time).
 - Moreover, all data read at any time must be committed; *i.e.*, REPEATABLE READ is a strictly stronger condition than READ COMMITTED .
 - Example: If 1 is executed before 3, then 2 must see the Bud and Miller tuples when it computes the min, even if it executes after 3. But if 1 executes between 3 and 4, then 2 may see the Heineken tuple.

- ❖ $X = \text{READ UNCOMMITTED}$: essentially no constraint, even on reading data written and then removed by a rollback.
 - Example: 1 and 2 could see Heineken, even if Joe rolled back his transaction.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Independence of Isolation Levels

Isolation levels describe what a transaction T with that isolation level sees.

- ❖ They *do not* constrain what other transactions, perhaps at different isolation levels, can see of the work done by T .

Example

If transaction 3-4 (Joe) runs serializable, but transaction 1-2 (Sally) does not, then Sally might see NULL as the value for both min and max, since it could appear to Sally that her transaction ran between steps 3 and 4.

Authorization in SQL

- ❖ File systems identify certain access privileges on files, *e.g.*, read, write, execute.
- ❖ In partial analogy, SQL identifies six access privileges on relations, of which the most important are:
 1. SELECT = the right to query the relation.
 2. INSERT = the right to insert tuples into the relation – may refer to one attribute, in which case the privilege is to specify only one column of the inserted tuple.
 3. DELETE = the right to delete tuples from the relation.
 4. UPDATE = the right to update tuples of the relation – may refer to one attribute.

Granting Privileges

- ❖ You have all possible privileges to the relations you create.
- ❖ You may grant privileges to any user if you have those privileges “with grant option.”
 - You have this option to your own relations.

Example

Assignment Project Exam Help

1. Here, Sally can query Sells and can change prices, but cannot pass on this power:

```
GRANT SELECT ON Sells,  
      UPDATE(price) ON Sells  
TO sally;
```

2. Here, Sally can also pass these privileges to whom she chooses:

```
GRANT SELECT ON Sells,  
      UPDATE(price) ON Sells  
TO sally  
WITH GRANT OPTION;
```


Revoking Privileges

- ❖ Your privileges can be revoked.
- ❖ Syntax is like granting, but REVOKE . . . FROM instead of GRANT . . . TO.
- ❖ Determining whether or not you have a privilege is tricky, involving “grant diagrams” as in text. However, the basic principles are:
 - a) If you have been given a privilege by several different people, then all of them have to revoke in order for you to lose the privilege.
 - b) Revocation is transitive. if A granted P to B , who then granted P to C , and then A revokes P from B , it is as if B also revoked P from C .

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder